

# BNN224

June 29, 2025

## 1 Binary Neural Network (BNN) for Plant Disease Classification

This notebook implements a Binary Neural Network using PyTorch for multiclass plant disease classification. The BNN uses binary weights and activations to reduce model size and computational requirements while maintaining reasonable accuracy.

### 1.1 Features:

- Binary weights and activations using sign function
- Processes 224x224 RGB images ( $3 \times 224 \times 224$  input)
- One hidden layer with binary weights
- Multiclass output with softmax activation
- CrossEntropyLoss for training

```
[4]: # Import Required Libraries
%pip install torch torchvision matplotlib numpy pandas seaborn scikit-learn
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset
import numpy as np
import matplotlib.pyplot as plt
from torchvision import transforms
import random

# Set random seeds for reproducibility
torch.manual_seed(42)
np.random.seed(42)
random.seed(42)

# Check if CUDA is available
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
print(f"Using device: {device}")
```

Requirement already satisfied: torch in ./venv/lib/python3.11/site-packages (2.7.1)

Requirement already satisfied: torchvision in ./venv/lib/python3.11/site-packages (0.22.1)

Requirement already satisfied: matplotlib in ./venv/lib/python3.11/site-packages (3.10.3)  
Requirement already satisfied: numpy in ./venv/lib/python3.11/site-packages (2.3.1)  
Requirement already satisfied: pandas in ./venv/lib/python3.11/site-packages (2.3.0)  
Requirement already satisfied: seaborn in ./venv/lib/python3.11/site-packages (0.13.2)  
Requirement already satisfied: scikit-learn in ./venv/lib/python3.11/site-packages (1.7.0)  
Requirement already satisfied: filelock in ./venv/lib/python3.11/site-packages (from torch) (3.18.0)  
Requirement already satisfied: typing-extensions>=4.10.0 in ./venv/lib/python3.11/site-packages (from torch) (4.14.0)  
Requirement already satisfied: sympy>=1.13.3 in ./venv/lib/python3.11/site-packages (from torch) (1.14.0)  
Requirement already satisfied: networkx in ./venv/lib/python3.11/site-packages (from torch) (3.5)  
Requirement already satisfied: jinja2 in ./venv/lib/python3.11/site-packages (from torch) (3.1.6)  
Requirement already satisfied: fsspec in ./venv/lib/python3.11/site-packages (from torch) (2025.5.1)  
Requirement already satisfied: nvidia-cuda-nvrtc-cu12==12.6.77 in ./venv/lib/python3.11/site-packages (from torch) (12.6.77)  
Requirement already satisfied: nvidia-cuda-runtime-cu12==12.6.77 in ./venv/lib/python3.11/site-packages (from torch) (12.6.77)  
Requirement already satisfied: nvidia-cuda-cupti-cu12==12.6.80 in ./venv/lib/python3.11/site-packages (from torch) (12.6.80)  
Requirement already satisfied: nvidia-cudnn-cu12==9.5.1.17 in ./venv/lib/python3.11/site-packages (from torch) (9.5.1.17)  
Requirement already satisfied: nvidia-cublas-cu12==12.6.4.1 in ./venv/lib/python3.11/site-packages (from torch) (12.6.4.1)  
Requirement already satisfied: nvidia-cufft-cu12==11.3.0.4 in ./venv/lib/python3.11/site-packages (from torch) (11.3.0.4)  
Requirement already satisfied: nvidia-curand-cu12==10.3.7.77 in ./venv/lib/python3.11/site-packages (from torch) (10.3.7.77)  
Requirement already satisfied: nvidia-cusolver-cu12==11.7.1.2 in ./venv/lib/python3.11/site-packages (from torch) (11.7.1.2)  
Requirement already satisfied: nvidia-cuspars-cu12==12.5.4.2 in ./venv/lib/python3.11/site-packages (from torch) (12.5.4.2)  
Requirement already satisfied: nvidia-cusparselt-cu12==0.6.3 in ./venv/lib/python3.11/site-packages (from torch) (0.6.3)  
Requirement already satisfied: nvidia-nccl-cu12==2.26.2 in ./venv/lib/python3.11/site-packages (from torch) (2.26.2)  
Requirement already satisfied: nvidia-nvtx-cu12==12.6.77 in ./venv/lib/python3.11/site-packages (from torch) (12.6.77)  
Requirement already satisfied: nvidia-nvjitlink-cu12==12.6.85 in ./venv/lib/python3.11/site-packages (from torch) (12.6.85)

Requirement already satisfied: nvidia-cufile-cu12==1.11.1.6 in  
 ./venv/lib/python3.11/site-packages (from torch) (1.11.1.6)  
 Requirement already satisfied: triton==3.3.1 in ./venv/lib/python3.11/site-  
 packages (from torch) (3.3.1)  
 Requirement already satisfied: setuptools>=40.8.0 in  
 ./venv/lib/python3.11/site-packages (from triton==3.3.1->torch) (80.9.0)  
 Requirement already satisfied: pillow!=8.3.\*,>=5.3.0 in  
 ./venv/lib/python3.11/site-packages (from torchvision) (11.2.1)  
 Requirement already satisfied: contourpy>=1.0.1 in ./venv/lib/python3.11/site-  
 packages (from matplotlib) (1.3.2)  
 Requirement already satisfied: cyclor>=0.10 in ./venv/lib/python3.11/site-  
 packages (from matplotlib) (0.12.1)  
 Requirement already satisfied: fonttools>=4.22.0 in ./venv/lib/python3.11/site-  
 packages (from matplotlib) (4.58.4)  
 Requirement already satisfied: kiwisolver>=1.3.1 in ./venv/lib/python3.11/site-  
 packages (from matplotlib) (1.4.8)  
 Requirement already satisfied: packaging>=20.0 in ./venv/lib/python3.11/site-  
 packages (from matplotlib) (25.0)  
 Requirement already satisfied: pyparsing>=2.3.1 in ./venv/lib/python3.11/site-  
 packages (from matplotlib) (3.2.3)  
 Requirement already satisfied: python-dateutil>=2.7 in  
 ./venv/lib/python3.11/site-packages (from matplotlib) (2.9.0.post0)  
 Requirement already satisfied: pytz>=2020.1 in ./venv/lib/python3.11/site-  
 packages (from pandas) (2025.2)  
 Requirement already satisfied: tzdata>=2022.7 in ./venv/lib/python3.11/site-  
 packages (from pandas) (2025.2)  
 Requirement already satisfied: scipy>=1.8.0 in ./venv/lib/python3.11/site-  
 packages (from scikit-learn) (1.16.0)  
 Requirement already satisfied: joblib>=1.2.0 in ./venv/lib/python3.11/site-  
 packages (from scikit-learn) (1.5.1)  
 Requirement already satisfied: threadpoolctl>=3.1.0 in  
 ./venv/lib/python3.11/site-packages (from scikit-learn) (3.6.0)  
 Requirement already satisfied: six>=1.5 in ./venv/lib/python3.11/site-packages  
 (from python-dateutil>=2.7->matplotlib) (1.17.0)  
 Requirement already satisfied: mpmath<1.4,>=1.1.0 in  
 ./venv/lib/python3.11/site-packages (from sympy>=1.13.3->torch) (1.3.0)  
 Requirement already satisfied: MarkupSafe>=2.0 in ./venv/lib/python3.11/site-  
 packages (from jinja2->torch) (3.0.2)  
 Note: you may need to restart the kernel to use updated packages.  
 Using device: cuda

```
[5]: # Additional imports for enhanced visualization and data export
import pandas as pd
import seaborn as sns
from sklearn.metrics import classification_report, confusion_matrix,
    accuracy_score
from sklearn.metrics import precision_recall_fscore_support
```

```

import time
import datetime
import os

# Set style for better plots
plt.style.use('seaborn-v0_8')
sns.set_palette("husl")

# Create results directory
os.makedirs('results', exist_ok=True)
print("Results directory created: ./results/")

```

Results directory created: ./results/

```

[6]: # Binary Activation Function
class BinaryActivation(torch.autograd.Function):
    """
    Binary activation function using the sign function.
    Forward:  $\text{sign}(x) = \{-1 \text{ if } x < 0, +1 \text{ if } x \geq 0\}$ 
    Backward: Straight-through estimator (STE) - passes gradients through
    ↪ unchanged
    """

    @staticmethod
    def forward(ctx, input):
        # Apply sign function: -1 for negative, +1 for non-negative
        return torch.sign(input)

    @staticmethod
    def backward(ctx, grad_output):
        # Straight-through estimator: pass gradients through unchanged
        # This allows gradients to flow back during training
        return grad_output

def binary_activation(x):
    """Wrapper function for binary activation"""
    return BinaryActivation.apply(x)

```

```

[7]: # Binary Linear Layer
class BinaryLinear(nn.Module):
    """
    Binary Linear layer with binary weights.
    Weights are binarized using the sign function during forward pass.
    """

    def __init__(self, in_features, out_features, bias=True):
        super(BinaryLinear, self).__init__()

```

```

        self.in_features = in_features
        self.out_features = out_features

        # Initialize weights using normal distribution
        self.weight = nn.Parameter(torch.randn(out_features, in_features) * 0.1)

        if bias:
            self.bias = nn.Parameter(torch.zeros(out_features))
        else:
            self.register_parameter('bias', None)

    def forward(self, input):
        # Binarize weights using sign function
        binary_weight = torch.sign(self.weight)

        # Perform linear transformation with binary weights
        output = F.linear(input, binary_weight, self.bias)

        return output

    def extra_repr(self):
        return f'in_features={self.in_features}, out_features={self.out_features}, bias={self.bias is not None}'

```

```

[8]: # Binary Neural Network Model
class BinaryNeuralNetwork(nn.Module):
    """
    Binary Neural Network for multiclass plant disease classification.

    Architecture:
    - Input: Flattened 224x224 RGB images (3*224*224 = 150528 features)
    - Hidden Layer: Binary linear layer with binary activation
    - Output Layer: Regular linear layer for class logits
    - Final: Softmax for multiclass prediction
    """

    def __init__(self, input_size=3*224*224, hidden_size=512, num_classes=3):
        super(BinaryNeuralNetwork, self).__init__()

        self.input_size = input_size
        self.hidden_size = hidden_size
        self.num_classes = num_classes

        # First layer: Regular linear layer (input preprocessing)
        self.input_layer = nn.Linear(input_size, hidden_size)

        # Hidden layer: Binary linear layer

```

```

self.hidden_layer = BinaryLinear(hidden_size, hidden_size)

# Output layer: Regular linear layer for final classification
self.output_layer = nn.Linear(hidden_size, num_classes)

# Dropout for regularization
self.dropout = nn.Dropout(0.2)

def forward(self, x):
    # Flatten input if it's not already flattened
    if len(x.shape) > 2:
        x = x.view(x.size(0), -1) # Flatten to (batch_size, input_size)

    # Input layer with ReLU activation
    x = F.relu(self.input_layer(x))
    x = self.dropout(x)

    # Hidden layer with binary weights and binary activation
    x = self.hidden_layer(x)
    x = binary_activation(x) # Binary activation function
    x = self.dropout(x)

    # Output layer (no activation - raw logits)
    logits = self.output_layer(x)

    return logits

def predict_proba(self, x):
    """Get class probabilities using softmax"""
    with torch.no_grad():
        logits = self.forward(x)
        probabilities = F.softmax(logits, dim=1)
    return probabilities

def predict(self, x):
    """Get predicted class labels"""
    with torch.no_grad():
        logits = self.forward(x)
        predictions = torch.argmax(logits, dim=1)
    return predictions

```

```

[9]: # Enhanced Training and Evaluation Functions
def train_bnn(model, train_loader, criterion, optimizer, device, num_epochs=10):
    """
    Enhanced training function with detailed metrics collection
    """
    model.train()

```

```

# Initialize tracking lists
train_losses = []
train_accuracies = []
epoch_times = []
learning_rates = []

# Detailed metrics per epoch
training_history = []

print(f"Starting training for {num_epochs} epochs...")
print("-" * 60)

start_time = time.time()

for epoch in range(num_epochs):
    epoch_start = time.time()

    running_loss = 0.0
    correct_predictions = 0
    total_samples = 0
    batch_losses = []

    # Training loop
    for batch_idx, (data, targets) in enumerate(train_loader):
        data, targets = data.to(device), targets.to(device)

        # Zero gradients
        optimizer.zero_grad()

        # Forward pass
        outputs = model(data)
        loss = criterion(outputs, targets)

        # Backward pass
        loss.backward()
        optimizer.step()

        # Statistics
        batch_loss = loss.item()
        running_loss += batch_loss
        batch_losses.append(batch_loss)

        _, predicted = torch.max(outputs.data, 1)
        total_samples += targets.size(0)
        correct_predictions += (predicted == targets).sum().item()

```

```

    # Calculate epoch metrics
    epoch_loss = running_loss / len(train_loader)
    epoch_accuracy = 100 * correct_predictions / total_samples
    epoch_time = time.time() - epoch_start
    current_lr = optimizer.param_groups[0]['lr']

    # Store metrics
    train_losses.append(epoch_loss)
    train_accuracies.append(epoch_accuracy)
    epoch_times.append(epoch_time)
    learning_rates.append(current_lr)

    # Store detailed history
    training_history.append({
        'epoch': epoch + 1,
        'loss': epoch_loss,
        'accuracy': epoch_accuracy,
        'time': epoch_time,
        'learning_rate': current_lr,
        'min_batch_loss': min(batch_losses),
        'max_batch_loss': max(batch_losses),
        'std_batch_loss': np.std(batch_losses)
    })

    # Progress display
    print(f'Epoch [{epoch+1:2d}/{num_epochs}] | Loss: {epoch_loss:.4f} | '
          f'Acc: {epoch_accuracy:6.2f}% | Time: {epoch_time:.2f}s | LR:␣
↪{current_lr:.6f}')

    total_time = time.time() - start_time
    print("-" * 60)
    print(f"Training completed in {total_time:.2f} seconds")
    print(f"Average epoch time: {np.mean(epoch_times):.2f}s")

    return train_losses, train_accuracies, training_history, epoch_times

def evaluate_bnn(model, test_loader, criterion, device, class_names):
    """
    Enhanced evaluation function with detailed metrics
    """
    model.eval()
    test_loss = 0.0
    correct_predictions = 0
    total_samples = 0

    all_predictions = []
    all_targets = []

```



```

all_probabilities = []

with torch.no_grad():
    for data, targets in test_loader:
        data, targets = data.to(device), targets.to(device)

        outputs = model(data)
        loss = criterion(outputs, targets)

        test_loss += loss.item()
        probabilities = F.softmax(outputs, dim=1)
        _, predicted = torch.max(outputs.data, 1)

        total_samples += targets.size(0)
        correct_predictions += (predicted == targets).sum().item()

        # Store for detailed analysis
        all_predictions.extend(predicted.cpu().numpy())
        all_targets.extend(targets.cpu().numpy())
        all_probabilities.extend(probabilities.cpu().numpy())

test_loss /= len(test_loader)
test_accuracy = 100 * correct_predictions / total_samples

print(f'Test Loss: {test_loss:.4f}, Test Accuracy: {test_accuracy:.2f}%')

# Generate detailed classification report
report = classification_report(all_targets, all_predictions,
                              target_names=class_names, output_dict=True)

# Calculate per-class metrics
precision, recall, fscore, support = precision_recall_fscore_support(
    all_targets, all_predictions, average=None,
    labels=range(len(class_names))
)

# Create detailed metrics dictionary
detailed_metrics = {
    'test_loss': test_loss,
    'test_accuracy': test_accuracy,
    'predictions': all_predictions,
    'targets': all_targets,
    'probabilities': all_probabilities,
    'classification_report': report,
    'per_class_precision': precision,
    'per_class_recall': recall,
    'per_class_fscore': fscore,

```

```

        'per_class_support': support
    }

    return detailed_metrics

```

```

[10]: # Load Real Plant Disease Dataset
from torchvision import datasets, transforms
from PIL import Image
import os

def load_plant_disease_dataset(dataset_path, image_size=224):
    """
    Load the real plant disease dataset from the specified path.
    Dataset structure:
    - dataset_path/Healthy_Soyabean/
    - dataset_path/rust/
    - dataset_path/Soyabean_Mosaic/
    """

    # Define transforms for preprocessing
    transform = transforms.Compose([
        transforms.Resize((image_size, image_size)), # Resize to 224x224
        transforms.ToTensor(), # Convert to tensor and normalize to [0,1]
        transforms.Normalize(mean=[0.485, 0.456, 0.406],
                              std=[0.229, 0.224, 0.225]) # ImageNet normalization
    ])

    # Load dataset using ImageFolder
    dataset = datasets.ImageFolder(root=dataset_path, transform=transform)

    # Convert to tensors
    data_loader = torch.utils.data.DataLoader(dataset, batch_size=len(dataset),
    ↪shuffle=False)
    X, y = next(iter(data_loader))

    # Print class mapping
    print("Class mapping:")
    for idx, class_name in enumerate(dataset.classes):
        print(f" {idx}: {class_name}")

    return X, y, dataset.classes

# Load your real plant disease dataset
dataset_path = "/home/dragon/Downloads/MH-SoyaHealthVision An Indian UAV and
    ↪Leaf Image Dataset for Integrated Crop Health Assessment/
    ↪Soyabean_UAV-Based_Image_Dataset"

```

```

# dataset_path = "/home/dragoon/Downloads/MH-SoyaHealthVision/
↳Soyabean_UAV-Based_Image_Dataset"
print("Loading real plant disease dataset...")
print(f"Dataset path: {dataset_path}")

X, y, class_names = load_plant_disease_dataset(dataset_path, image_size=224)

print(f"\nDataset loaded successfully!")
print(f"Dataset shape: {X.shape}")
print(f"Labels shape: {y.shape}")
print(f"Number of classes: {len(torch.unique(y))}")
print(f"Class distribution: {torch.bincount(y)}")
print(f"Class names: {class_names}")

# Display some dataset statistics
print(f"\nDataset Statistics:")
for i, class_name in enumerate(class_names):
    count = (y == i).sum().item()
    print(f"  {class_name}: {count} images")

```

Loading real plant disease dataset...

Dataset path: /home/dragoon/Downloads/MH-SoyaHealthVision An Indian UAV and Leaf Image Dataset for Integrated Crop Health Assessment/Soyabean\_UAV-Based\_Image\_Dataset

Class mapping:

```

0: Healthy_Soyabean
1: Soyabean_Semilooper_Pest_Attack
2: Soyabean_Mosaic
3: rust

```

Dataset loaded successfully!

Dataset shape: torch.Size([2842, 3, 224, 224])

Labels shape: torch.Size([2842])

Number of classes: 4

Class distribution: tensor([ 280, 790, 772, 1000])

Class names: ['Healthy\_Soyabean', 'Soyabean\_Semilooper\_Pest\_Attack', 'Soyabean\_Mosaic', 'rust']

Dataset Statistics:

```

Healthy_Soyabean: 280 images
Soyabean_Semilooper_Pest_Attack: 790 images
Soyabean_Mosaic: 772 images
rust: 1000 images

```

```

[11]: # Split Dataset and Create Data Loaders
      %pip install scikit-learn
      from sklearn.model_selection import train_test_split

```

```

# Split dataset into train and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)

print(f"Training set: {X_train.shape[0]} samples")
print(f"Test set: {X_test.shape[0]} samples")

# Create data loaders
batch_size = 32

train_dataset = TensorDataset(X_train, y_train)
test_dataset = TensorDataset(X_test, y_test)

train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=False)

print(f"Number of training batches: {len(train_loader)}")
print(f"Number of test batches: {len(test_loader)}")

```

Requirement already satisfied: scikit-learn in ./venv/lib/python3.11/site-packages (1.7.0)

Requirement already satisfied: numpy>=1.22.0 in ./venv/lib/python3.11/site-packages (from scikit-learn) (2.3.1)

Requirement already satisfied: scipy>=1.8.0 in ./venv/lib/python3.11/site-packages (from scikit-learn) (1.16.0)

Requirement already satisfied: joblib>=1.2.0 in ./venv/lib/python3.11/site-packages (from scikit-learn) (1.5.1)

Requirement already satisfied: threadpoolctl>=3.1.0 in ./venv/lib/python3.11/site-packages (from scikit-learn) (3.6.0)

Note: you may need to restart the kernel to use updated packages.

Training set: 2273 samples

Test set: 569 samples

Number of training batches: 72

Number of test batches: 18

```

[12]: # Visualize Sample Images from the Dataset
def denormalize_image(tensor, mean=[0.485, 0.456, 0.406], std=[0.229, 0.224, 0.
    ↪225]):
    """Denormalize image tensor for visualization"""
    mean = torch.tensor(mean).view(3, 1, 1)
    std = torch.tensor(std).view(3, 1, 1)
    return tensor * std + mean

# Display sample images from each class
plt.figure(figsize=(15, 5))

```

```

samples_per_class = 3
for class_idx, class_name in enumerate(class_names):
    # Get indices for this class
    class_indices = torch.where(y == class_idx)[0]

    # Select random samples from this class
    random_indices = torch.randperm(len(class_indices))[:samples_per_class]
    sample_indices = class_indices[random_indices]

    for i, sample_idx in enumerate(sample_indices):
        plt.subplot(len(class_names), samples_per_class, class_idx *
↪samples_per_class + i + 1)

        # Get and denormalize image
        img = denormalize_image(X[sample_idx])
        img = torch.clamp(img, 0, 1) # Ensure values are in [0,1]


        # Convert to numpy and transpose for matplotlib
        img_np = img.permute(1, 2, 0).numpy()

        plt.imshow(img_np)
        plt.title(f"{class_name}\n(Class {class_idx})")
        plt.axis('off')

plt.tight_layout()
plt.suptitle('Sample Images from Plant Disease Dataset', y=1.02, fontsize=16)
plt.show()

print(f"Sample images from your plant disease dataset:")
print(f"Classes: {'', '.join(class_names)}")
print(f"Image size: 224x224 pixels")
print(f"Total images: {len(X)}")

```



Sample images from your plant disease dataset:

Classes: Healthy\_Soyabean, Soyabean Semilooper\_Pest\_Attack, Soyabean\_Mosaic, rust

Image size: 224x224 pixels

Total images: 2842

```
[13]: # Initialize and Train the BNN Model
# Model parameters
input_size = 3 * 224 * 224 # 224x224 RGB images (increased resolution)
hidden_size = 512
num_classes = len(class_names) # Dynamic based on actual dataset

print(f"Model Configuration:")
print(f"  Input size: {input_size}")
print(f"  Hidden size: {hidden_size}")
print(f"  Number of classes: {num_classes}")
print(f"  Classes: {class_names}")

# Initialize model
model = BinaryNeuralNetwork(
    input_size=input_size,
    hidden_size=hidden_size,
    num_classes=num_classes
).to(device)

# Print model architecture
print("\nBinary Neural Network Architecture:")
print(model)
print(f"\nTotal parameters: {sum(p.numel() for p in model.parameters()):,}")
```

```

print(f"Trainable parameters: {sum(p.numel() for p in model.parameters() if p.
    ↪requires_grad):,}")

# Define loss function and optimizer
criterion = nn.CrossEntropyLoss()
optimizer = optim.Adam(model.parameters(), lr=0.001, weight_decay=1e-4)

# Train the model with enhanced tracking
print("\nStarting enhanced training with detailed metrics...")
train_losses, train_accuracies, training_history, epoch_times = train_bnn(
    model=model,
    train_loader=train_loader,
    criterion=criterion,
    optimizer=optimizer,
    device=device,
    num_epochs=20
)

```

Model Configuration:

Input size: 150528  
Hidden size: 512  
Number of classes: 4  
Classes: ['Healthy\_Soyabean', 'Soyabean Semilooper\_Pest\_Attack',  
'Soyabean\_Mosaic', 'rust']

Binary Neural Network Architecture:

```

BinaryNeuralNetwork(
  (input_layer): Linear(in_features=150528, out_features=512, bias=True)
  (hidden_layer): BinaryLinear(in_features=512, out_features=512, bias=True)
  (output_layer): Linear(in_features=512, out_features=4, bias=True)
  (dropout): Dropout(p=0.2, inplace=False)
)

```

Total parameters: 77,335,556

Trainable parameters: 77,335,556

Starting enhanced training with detailed metrics...

Starting training for 20 epochs...

```

-----
Epoch [ 1/20] | Loss: 1.0177 | Acc: 56.58% | Time: 3.72s | LR: 0.001000
Epoch [ 2/20] | Loss: 0.8065 | Acc: 64.32% | Time: 3.49s | LR: 0.001000
Epoch [ 3/20] | Loss: 0.7218 | Acc: 69.69% | Time: 3.49s | LR: 0.001000
Epoch [ 4/20] | Loss: 0.6995 | Acc: 68.50% | Time: 3.49s | LR: 0.001000
Epoch [ 5/20] | Loss: 0.6289 | Acc: 72.94% | Time: 3.48s | LR: 0.001000
Epoch [ 6/20] | Loss: 0.6023 | Acc: 74.88% | Time: 3.49s | LR: 0.001000
Epoch [ 7/20] | Loss: 0.6255 | Acc: 72.68% | Time: 3.49s | LR: 0.001000
Epoch [ 8/20] | Loss: 0.6125 | Acc: 74.44% | Time: 3.48s | LR: 0.001000
Epoch [ 9/20] | Loss: 0.6179 | Acc: 74.66% | Time: 3.47s | LR: 0.001000

```

Epoch [10/20]	Loss: 0.6904	Acc: 71.01%	Time: 3.48s	LR: 0.001000
Epoch [11/20]	Loss: 0.6373	Acc: 72.20%	Time: 3.47s	LR: 0.001000
Epoch [12/20]	Loss: 0.6937	Acc: 71.54%	Time: 3.47s	LR: 0.001000
Epoch [13/20]	Loss: 0.7616	Acc: 66.78%	Time: 3.46s	LR: 0.001000
Epoch [14/20]	Loss: 0.6925	Acc: 70.66%	Time: 3.47s	LR: 0.001000
Epoch [15/20]	Loss: 0.7682	Acc: 65.24%	Time: 3.47s	LR: 0.001000
Epoch [16/20]	Loss: 0.8242	Acc: 63.97%	Time: 3.47s	LR: 0.001000
Epoch [17/20]	Loss: 0.9051	Acc: 59.17%	Time: 3.48s	LR: 0.001000
Epoch [18/20]	Loss: 0.9313	Acc: 59.04%	Time: 3.47s	LR: 0.001000
Epoch [19/20]	Loss: 1.0035	Acc: 57.11%	Time: 3.47s	LR: 0.001000
Epoch [20/20]	Loss: 1.0514	Acc: 54.73%	Time: 3.47s	LR: 0.001000

-----

Training completed in 69.82 seconds

Average epoch time: 3.49s

```
[14]: # Enhanced Model Evaluation and Comprehensive Visualization
print("\nEvaluating model with detailed metrics...")
detailed_metrics = evaluate_bnn(model, test_loader, criterion, device,
    ↪class_names)

# Create comprehensive visualizations
fig = plt.figure(figsize=(20, 15))

# 1. Training Loss and Accuracy
plt.subplot(3, 4, 1)
plt.plot(range(1, len(train_losses) + 1), train_losses, 'b-', linewidth=2,
    ↪label='Training Loss')
plt.title('Training Loss Over Epochs', fontsize=12, fontweight='bold')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid(True, alpha=0.3)
plt.legend()

plt.subplot(3, 4, 2)
plt.plot(range(1, len(train_accuracies) + 1), train_accuracies, 'r-',
    ↪linewidth=2, label='Training Accuracy')
plt.title('Training Accuracy Over Epochs', fontsize=12, fontweight='bold')
plt.xlabel('Epoch')
plt.ylabel('Accuracy (%)')
plt.grid(True, alpha=0.3)
plt.legend()

# 2. Training Time Analysis
plt.subplot(3, 4, 3)
plt.bar(range(1, len(epoch_times) + 1), epoch_times, alpha=0.7, color='green')
plt.title('Training Time per Epoch', fontsize=12, fontweight='bold')
plt.xlabel('Epoch')
```



```

plt.ylabel('Time (seconds)')
plt.grid(True, alpha=0.3)

# 3. Confusion Matrix
plt.subplot(3, 4, 4)
cm = confusion_matrix(detailed_metrics['targets'],
    ↪detailed_metrics['predictions'])
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
    xticklabels=class_names, yticklabels=class_names)
plt.title('Confusion Matrix', fontsize=12, fontweight='bold')
plt.ylabel('True Label')
plt.xlabel('Predicted Label')

# 4. Per-Class Performance
plt.subplot(3, 4, 5)
metrics_df = pd.DataFrame({
    'Precision': detailed_metrics['per_class_precision'],
    'Recall': detailed_metrics['per_class_recall'],
    'F1-Score': detailed_metrics['per_class_fscore']
}, index=class_names)

metrics_df.plot(kind='bar', ax=plt.gca(), width=0.8)
plt.title('Per-Class Performance Metrics', fontsize=12, fontweight='bold')
plt.ylabel('Score')
plt.legend(loc='upper right')
plt.xticks(rotation=45)

# 5. Class Distribution
plt.subplot(3, 4, 6)
class_counts = [sum(1 for label in detailed_metrics['targets'] if label == i)
    ↪for i in range(len(class_names))]
plt.pie(class_counts, labels=class_names, autopct='%1.1f%%', startangle=90)
plt.title('Test Set Class Distribution', fontsize=12, fontweight='bold')

# 6. Loss Smoothness (Moving Average)
plt.subplot(3, 4, 7)
window_size = 5
if len(train_losses) >= window_size:
    smoothed_loss = pd.Series(train_losses).rolling(window=window_size).mean()
    plt.plot(range(1, len(train_losses) + 1), train_losses, 'b-', alpha=0.3,
    ↪label='Raw Loss')
    plt.plot(range(1, len(smoothed_loss) + 1), smoothed_loss, 'b-',
    ↪linewidth=2, label=f'{window_size}-Epoch Moving Avg')
    plt.legend()
else:
    plt.plot(range(1, len(train_losses) + 1), train_losses, 'b-', linewidth=2,
    ↪label='Training Loss')

```

```

plt.title('Loss Smoothness Analysis', fontsize=12, fontweight='bold')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.grid(True, alpha=0.3)

# 7. Prediction Confidence Distribution
plt.subplot(3, 4, 8)
max_probs = [max(prob) for prob in detailed_metrics['probabilities']]
plt.hist(max_probs, bins=20, alpha=0.7, color='purple', edgecolor='black')
plt.title('Prediction Confidence Distribution', fontsize=12, fontweight='bold')
plt.xlabel('Maximum Probability')
plt.ylabel('Frequency')
plt.grid(True, alpha=0.3)

# 8. Learning Rate Over Time (if scheduler was used)
plt.subplot(3, 4, 9)
learning_rates = [history['learning_rate'] for history in training_history]
plt.plot(range(1, len(learning_rates) + 1), learning_rates, 'orange',
         ↪ linewidth=2)
plt.title('Learning Rate Schedule', fontsize=12, fontweight='bold')
plt.xlabel('Epoch')
plt.ylabel('Learning Rate')
plt.grid(True, alpha=0.3)

# 9. Training Metrics Summary
plt.subplot(3, 4, 10)
summary_metrics = [
    detailed_metrics['test_accuracy'],
    np.mean(detailed_metrics['per_class_precision']) * 100,
    np.mean(detailed_metrics['per_class_recall']) * 100,
    np.mean(detailed_metrics['per_class_fscore']) * 100
]
metric_names = ['Accuracy', 'Avg Precision', 'Avg Recall', 'Avg F1-Score']
bars = plt.bar(metric_names, summary_metrics, color=['blue', 'green', 'orange',
         ↪ 'red'], alpha=0.7)
plt.title('Model Performance Summary', fontsize=12, fontweight='bold')
plt.ylabel('Score (%)')
plt.ylim(0, 100)
# Add value labels on bars
for bar, value in zip(bars, summary_metrics):
    plt.text(bar.get_x() + bar.get_width()/2, bar.get_height() + 1,
             f'{value:.1f}%', ha='center', va='bottom', fontweight='bold')
plt.xticks(rotation=45)

# 10. Error Analysis
plt.subplot(3, 4, 11)

```

```

correct_mask = np.array(detailed_metrics['targets']) == np.
    ↳array(detailed_metrics['predictions'])
correct_confidences = [max(prob) for i, prob in
    ↳enumerate(detailed_metrics['probabilities']) if correct_mask[i]]
incorrect_confidences = [max(prob) for i, prob in
    ↳enumerate(detailed_metrics['probabilities']) if not correct_mask[i]]

plt.hist(correct_confidences, bins=15, alpha=0.7, label='Correct Predictions',
    ↳color='green')
plt.hist(incorrect_confidences, bins=15, alpha=0.7, label='Incorrect
    ↳Predictions', color='red')
plt.title('Confidence: Correct vs Incorrect', fontsize=12, fontweight='bold')
plt.xlabel('Prediction Confidence')
plt.ylabel('Frequency')
plt.legend()
plt.grid(True, alpha=0.3)

# 11. Loss Improvement Rate
plt.subplot(3, 4, 12)
loss_improvements = [-1 * (train_losses[i] - train_losses[i-1]) for i in
    ↳range(1, len(train_losses))]
plt.plot(range(2, len(train_losses) + 1), loss_improvements, 'purple',
    ↳marker='o', linewidth=2)
plt.title('Loss Improvement Rate', fontsize=12, fontweight='bold')
plt.xlabel('Epoch')
plt.ylabel('Loss Improvement')
plt.grid(True, alpha=0.3)
plt.axhline(y=0, color='black', linestyle='--', alpha=0.5)

plt.tight_layout()
plt.savefig('results/comprehensive_training_analysis.png', dpi=300,
    ↳bbox_inches='tight')
plt.show()

print(f"\n{'='*60}")
print("COMPREHENSIVE MODEL PERFORMANCE REPORT")
print(f"\n{'='*60}")
print(f"Overall Test Accuracy: {detailed_metrics['test_accuracy']:.2f}%")
print(f"Overall Test Loss: {detailed_metrics['test_loss']:.4f}")
print(f"Total Training Time: {sum(epoch_times):.2f} seconds")
print(f"Average Epoch Time: {np.mean(epoch_times):.2f} seconds")
print(f"\nPer-Class Performance:")
for i, class_name in enumerate(class_names):
    print(f"    {class_name}:")
    print(f"        Precision: {detailed_metrics['per_class_precision'][i]:.3f}")
    print(f"        Recall: {detailed_metrics['per_class_recall'][i]:.3f}")

```

```

    print(f"    F1-Score: {detailed_metrics['per_class_fscore'][i]:.3f}")
    print(f"    Support: {detailed_metrics['per_class_support'][i]}")
print(f"{'='*60}")

# Test model predictions on a few samples
print("\nTesting model predictions:")
model.eval()
with torch.no_grad():
    # Get a small batch from test set
    test_samples = X_test[:5].to(device)
    test_labels = y_test[:5]

    # Get predictions
    logits = model(test_samples)
    probabilities = F.softmax(logits, dim=1)
    predicted_classes = torch.argmax(logits, dim=1)

    print("\nSample predictions:")
    for i in range(5):
        true_class_name = class_names[test_labels[i].item()]
        pred_class_name = class_names[predicted_classes[i].item()]

        print(f"Sample {i+1}:")
        print(f"  True label: {test_labels[i].item()} ({true_class_name})")
        print(f"  Predicted: {predicted_classes[i].item()} ({pred_class_name})")
        print(f"  Probabilities: {probabilities[i].cpu().numpy()}")

    # Show probability for each class
    for j, class_name in enumerate(class_names):
        prob = probabilities[i][j].item()
        print(f"    {class_name}: {prob:.3f}")
    print()

```

Evaluating model with detailed metrics...

Test Loss: 0.9215, Test Accuracy: 66.08%

```

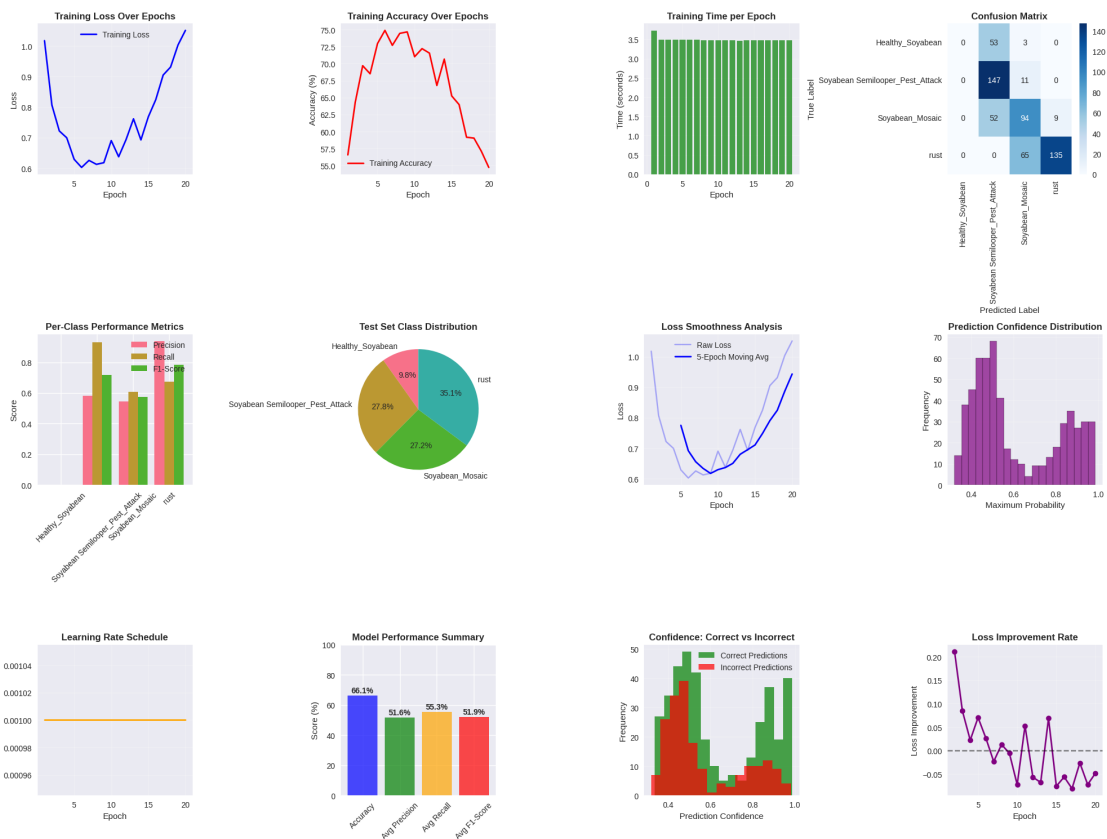
/home/dragoon/coding/drone-crop/BNN/binary nerural
network/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
/home/dragoon/coding/drone-crop/BNN/binary nerural
network/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.

```

```

_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
/home/dragoon/coding/drone-crop/BNN/binary nerural
network/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])
/home/dragoon/coding/drone-crop/BNN/binary nerural
network/.venv/lib/python3.11/site-
packages/sklearn/metrics/_classification.py:1706: UndefinedMetricWarning:
Precision is ill-defined and being set to 0.0 in labels with no predicted
samples. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, f"{metric.capitalize()} is", result.shape[0])

```



## COMPREHENSIVE MODEL PERFORMANCE REPORT

Overall Test Accuracy: 66.08%  
 Overall Test Loss: 0.9215  
 Total Training Time: 69.82 seconds

Average Epoch Time: 3.49 seconds

Per-Class Performance:

Healthy\_Soyabean:

Precision: 0.000

Recall: 0.000

F1-Score: 0.000

Support: 56

Soyabean\_Semilooper\_Pest\_Attack:

Precision: 0.583

Recall: 0.930

F1-Score: 0.717

Support: 158

Soyabean\_Mosaic:

Precision: 0.543

Recall: 0.606

F1-Score: 0.573

Support: 155

rust:

Precision: 0.938

Recall: 0.675

F1-Score: 0.785

Support: 200

=====

Testing model predictions:

Sample predictions:

Sample 1:

True label: 1 (Soyabean\_Semilooper\_Pest\_Attack)

Predicted: 1 (Soyabean\_Semilooper\_Pest\_Attack)

Probabilities: [0.08531532 0.8166505 0.06423339 0.03380084]

Healthy\_Soyabean: 0.085

Soyabean\_Semilooper\_Pest\_Attack: 0.817

Soyabean\_Mosaic: 0.064

rust: 0.034

Sample 2:

True label: 3 (rust)

Predicted: 3 (rust)

Probabilities: [0.07985441 0.12567203 0.26099566 0.53347784]

Healthy\_Soyabean: 0.080

Soyabean\_Semilooper\_Pest\_Attack: 0.126

Soyabean\_Mosaic: 0.261

rust: 0.533

Sample 3:

True label: 1 (Soyabean\_Semilooper\_Pest\_Attack)

Predicted: 1 (Soyabean Semilooper\_Pest\_Attack)  
Probabilities: [0.03808016 0.94920504 0.00672773 0.00598706]  
Healthy\_Soyabean: 0.038  
Soyabean Semilooper\_Pest\_Attack: 0.949  
Soyabean\_Mosaic: 0.007  
rust: 0.006

Sample 4:

True label: 2 (Soyabean\_Mosaic)  
Predicted: 1 (Soyabean Semilooper\_Pest\_Attack)  
Probabilities: [0.08994745 0.3865225 0.33603245 0.18749759]  
Healthy\_Soyabean: 0.090  
Soyabean Semilooper\_Pest\_Attack: 0.387  
Soyabean\_Mosaic: 0.336  
rust: 0.187

Sample 5:

True label: 3 (rust)  
Predicted: 2 (Soyabean\_Mosaic)  
Probabilities: [0.07532449 0.1616913 0.41420275 0.34878147]  
Healthy\_Soyabean: 0.075  
Soyabean Semilooper\_Pest\_Attack: 0.162  
Soyabean\_Mosaic: 0.414  
rust: 0.349

```
[15]: # Comprehensive CSV Data Export
print("\nExporting training data to CSV files...")

# 1. Training History CSV
training_df = pd.DataFrame(training_history)
training_df.to_csv('results/training_history.csv', index=False)
print(" Training history saved to: results/training_history.csv")

# 2. Detailed Test Results CSV
test_results = []
for i in range(len(detailed_metrics['targets'])):
    test_results.append({
        'sample_id': i,
        'true_label': detailed_metrics['targets'][i],
        'true_class': class_names[detailed_metrics['targets'][i]],
        'predicted_label': detailed_metrics['predictions'][i],
        'predicted_class': class_names[detailed_metrics['predictions'][i]],
        'correct': detailed_metrics['targets'][i] ==
↪detailed_metrics['predictions'][i],
        'confidence': max(detailed_metrics['probabilities'][i]),
```

```

        **{f'prob_{class_names[j]}': detailed_metrics['probabilities'][i][j]}
    for j in range(len(class_names)):
    })

test_results_df = pd.DataFrame(test_results)
test_results_df.to_csv('results/test_predictions.csv', index=False)
print(" Test predictions saved to: results/test_predictions.csv")

# 3. Per-Class Performance Metrics CSV
per_class_metrics = pd.DataFrame({
    'class_name': class_names,
    'precision': detailed_metrics['per_class_precision'],
    'recall': detailed_metrics['per_class_recall'],
    'f1_score': detailed_metrics['per_class_fscore'],
    'support': detailed_metrics['per_class_support']
})
per_class_metrics.to_csv('results/per_class_metrics.csv', index=False)
print(" Per-class metrics saved to: results/per_class_metrics.csv")

# 4. Model Configuration and Final Results CSV
model_summary = {
    'parameter': [
        'model_type', 'input_size', 'hidden_size', 'num_classes',
        'num_epochs', 'batch_size', 'learning_rate', 'optimizer',
        'total_parameters', 'trainable_parameters',
        'final_train_loss', 'final_train_accuracy', 'test_loss',
    test_accuracy',
        'total_training_time', 'avg_epoch_time',
        'dataset_total_samples', 'train_samples', 'test_samples'
    ],
    'value': [
        'Binary Neural Network', input_size, hidden_size, num_classes,
        20, batch_size, 0.001, 'Adam',
        sum(p.numel() for p in model.parameters()),
        sum(p.numel() for p in model.parameters() if p.requires_grad),
        train_losses[-1], train_accuracies[-1],
        detailed_metrics['test_loss'], detailed_metrics['test_accuracy'],
        sum(epoch_times), np.mean(epoch_times),
        len(X), len(X_train), len(X_test)
    ]
}

model_summary_df = pd.DataFrame(model_summary)
model_summary_df.to_csv('results/model_summary.csv', index=False)
print(" Model summary saved to: results/model_summary.csv")

# 5. Confusion Matrix CSV

```



```

cm = confusion_matrix(detailed_metrics['targets'],
    ↪detailed_metrics['predictions'])
cm_df = pd.DataFrame(cm, index=[f'True_{name}' for name in class_names],
    columns=[f'Pred_{name}' for name in class_names])
cm_df.to_csv('results/confusion_matrix.csv')
print(" Confusion matrix saved to: results/confusion_matrix.csv")

# 6. Training Progress Summary CSV
epoch_summary = []
for i in range(len(train_losses)):
    epoch_summary.append({
        'epoch': i + 1,
        'train_loss': train_losses[i],
        'train_accuracy': train_accuracies[i],
        'epoch_time': epoch_times[i],
        'cumulative_time': sum(epoch_times[:i+1]),
        'loss_improvement': 0 if i == 0 else train_losses[i-1] -
    ↪train_losses[i],
        'accuracy_improvement': 0 if i == 0 else train_accuracies[i] -
    ↪train_accuracies[i-1]
    })

epoch_summary_df = pd.DataFrame(epoch_summary)
epoch_summary_df.to_csv('results/epoch_summary.csv', index=False)
print(" Epoch summary saved to: results/epoch_summary.csv")

# 7. Binary Weights Analysis CSV
with torch.no_grad():
    hidden_weights = model.hidden_layer.weight.data.cpu().numpy()
    binary_weights = np.sign(hidden_weights)

    weights_analysis = {
        'layer': ['hidden_layer'],
        'total_weights': [hidden_weights.size],
        'positive_weights': [np.sum(binary_weights > 0)],
        'negative_weights': [np.sum(binary_weights < 0)],
        'zero_weights': [np.sum(binary_weights == 0)],
        'weight_mean': [np.mean(hidden_weights)],
        'weight_std': [np.std(hidden_weights)],
        'binary_weight_ratio': [np.sum(binary_weights > 0) / hidden_weights.
    ↪size]
    }

weights_analysis_df = pd.DataFrame(weights_analysis)
weights_analysis_df.to_csv('results/binary_weights_analysis.csv', index=False)
print(" Binary weights analysis saved to: results/binary_weights_analysis.csv")

```

```

# 8. Dataset Statistics CSV
dataset_stats = []
for i, class_name in enumerate(class_names):
    train_count = sum(1 for label in y_train if label == i)
    test_count = sum(1 for label in y_test if label == i)
    total_count = train_count + test_count

    dataset_stats.append({
        'class_name': class_name,
        'class_id': i,
        'train_samples': train_count,
        'test_samples': test_count,
        'total_samples': total_count,
        'train_percentage': (train_count / len(y_train)) * 100,
        'test_percentage': (test_count / len(y_test)) * 100,
        'overall_percentage': (total_count / len(y)) * 100
    })

dataset_stats_df = pd.DataFrame(dataset_stats)
dataset_stats_df.to_csv('results/dataset_statistics.csv', index=False)
print(" Dataset statistics saved to: results/dataset_statistics.csv")

# Create a comprehensive summary report
print(f"\n{'='*80}")
print("CSV FILES EXPORTED SUCCESSFULLY")
print(f"{'='*80}")
print("The following CSV files have been created in the 'results/' directory:")
print("1. training_history.csv - Detailed epoch-by-epoch training metrics")
print("2. test_predictions.csv - Individual test sample predictions and_
↳probabilities")
print("3. per_class_metrics.csv - Precision, recall, F1-score for each class")
print("4. model_summary.csv - Model configuration and final performance")
print("5. confusion_matrix.csv - Confusion matrix data")
print("6. epoch_summary.csv - Training progress with improvements")
print("7. binary_weights_analysis.csv - Analysis of binary weight distribution")
print("8. dataset_statistics.csv - Dataset composition and class distribution")
print(f"{'='*80}")

# Display sample of key CSV files
print("\nSample of Training History:")
print(training_df.head())
print("\nSample of Test Predictions:")
print(test_results_df.head())
print("\nPer-Class Metrics:")
print(per_class_metrics)

```

Exporting training data to CSV files...

Training history saved to: results/training\_history.csv  
 Test predictions saved to: results/test\_predictions.csv  
 Per-class metrics saved to: results/per\_class\_metrics.csv  
 Model summary saved to: results/model\_summary.csv  
 Confusion matrix saved to: results/confusion\_matrix.csv  
 Epoch summary saved to: results/epoch\_summary.csv  
 Binary weights analysis saved to: results/binary\_weights\_analysis.csv  
 Dataset statistics saved to: results/dataset\_statistics.csv

=====

CSV FILES EXPORTED SUCCESSFULLY

=====

The following CSV files have been created in the 'results/' directory:

1. training\_history.csv - Detailed epoch-by-epoch training metrics
  2. test\_predictions.csv - Individual test sample predictions and probabilities
  3. per\_class\_metrics.csv - Precision, recall, F1-score for each class
  4. model\_summary.csv - Model configuration and final performance
  5. confusion\_matrix.csv - Confusion matrix data
  6. epoch\_summary.csv - Training progress with improvements
  7. binary\_weights\_analysis.csv - Analysis of binary weight distribution
  8. dataset\_statistics.csv - Dataset composition and class distribution
- =====

Sample of Training History:

	epoch	loss	accuracy	time	learning_rate	min_batch_loss	\
0	1	1.017667	56.577211	3.723100	0.001	0.637834	
1	2	0.806459	64.320282	3.494545	0.001	0.471640	
2	3	0.721814	69.687637	3.491834	0.001	0.516664	
3	4	0.699459	68.499780	3.492792	0.001	0.090115	
4	5	0.628917	72.943247	3.484508	0.001	0.291184	

	max_batch_loss	std_batch_loss
0	2.485080	0.289303
1	1.157512	0.131777
2	1.849854	0.172884
3	1.158462	0.140914
4	1.007220	0.114069

Sample of Test Predictions:

	sample_id	true_label	true_class	predicted_label	\
0	0	1	Soyabean Semilooper_Pest_Attack	1	
1	1	3	rust	3	
2	2	1	Soyabean Semilooper_Pest_Attack	1	
3	3	2	Soyabean_Mosaic	1	
4	4	3	rust	2	

	predicted_class	correct	confidence	\
0	Soyabean Semilooper_Pest_Attack	True	0.816651	

1		rust	True	0.533478
2	Soyabean	Semilooper_Pest_Attack	True	0.949205
3	Soyabean	Semilooper_Pest_Attack	False	0.386523
4		Soyabean_Mosaic	False	0.414203

	prob_Healthy_Soyabean	prob_Soyabean Semilooper_Pest_Attack	\
0	0.085315		0.816651
1	0.079854		0.125672
2	0.038080		0.949205
3	0.089947		0.386523
4	0.075324		0.161691

	prob_Soyabean_Mosaic	prob_rust
0	0.064233	0.033801
1	0.260996	0.533478
2	0.006728	0.005987
3	0.336032	0.187498
4	0.414203	0.348781

Per-Class Metrics:

	class_name	precision	recall	f1_score	support
0	Healthy_Soyabean	0.000000	0.000000	0.000000	56
1	Soyabean Semilooper_Pest_Attack	0.583333	0.930380	0.717073	158
2	Soyabean_Mosaic	0.543353	0.606452	0.573171	155
3	rust	0.937500	0.675000	0.784884	200

```
[16]: # Model Saving and Additional Analysis
print("\nSaving model and creating additional analyses...")

# Save the trained model
timestamp = datetime.datetime.now().strftime("%Y%m%d_%H%M%S")
model_path = f'results/bnn_plant_disease_model_{timestamp}.pth'
torch.save(model.state_dict(), model_path)
print(f" Model saved to: {model_path}")

# Save model architecture info
model_info = {
    'model_class': 'BinaryNeuralNetwork',
    'input_size': input_size,
    'hidden_size': hidden_size,
    'num_classes': num_classes,
    'class_names': class_names,
    'timestamp': timestamp,
    'final_accuracy': detailed_metrics['test_accuracy'],
    'final_loss': detailed_metrics['test_loss']
}
```

```

import json
with open(f'results/model_info_{timestamp}.json', 'w') as f:
    json.dump(model_info, f, indent=2)
print(f" Model info saved to: results/model_info_{timestamp}.json")

# Create ROC Curves for each class (one-vs-rest)
from sklearn.metrics import roc_curve, auc
from sklearn.preprocessing import label_binarize

# Binarize the output for ROC calculation
y_test_bin = label_binarize(detailed_metrics['targets'],
    ↳classes=range(len(class_names)))
y_score = np.array(detailed_metrics['probabilities'])

plt.figure(figsize=(15, 5))

# Plot ROC curve for each class
for i in range(len(class_names)):
    plt.subplot(1, 3, i+1)
    fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc = auc(fpr, tpr)

    plt.plot(fpr, tpr, linewidth=2, label=f'ROC curve (AUC = {roc_auc:.3f})')
    plt.plot([0, 1], [0, 1], 'k--', linewidth=1)
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.05])
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')
    plt.title(f'ROC Curve - {class_names[i]}')
    plt.legend(loc="lower right")
    plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('results/roc_curves.png', dpi=300, bbox_inches='tight')
plt.show()

# Create improvement trends visualization
plt.figure(figsize=(16, 8))

# Loss improvement over time
plt.subplot(2, 2, 1)
loss_smooth = pd.Series(train_losses).rolling(window=3, center=True).mean()
plt.plot(range(1, len(train_losses) + 1), train_losses, 'b-', alpha=0.5,
    ↳label='Raw Loss')
plt.plot(range(1, len(loss_smooth) + 1), loss_smooth, 'b-', linewidth=2,
    ↳label='Smoothed Loss')
plt.title('Training Loss Improvement Trend')

```

```

plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True, alpha=0.3)

# Accuracy improvement over time
plt.subplot(2, 2, 2)
acc_smooth = pd.Series(train_accuracies).rolling(window=3, center=True).mean()
plt.plot(range(1, len(train_accuracies) + 1), train_accuracies, 'r-', alpha=0.
    ↳5, label='Raw Accuracy')
plt.plot(range(1, len(acc_smooth) + 1), acc_smooth, 'r-', linewidth=2,
    ↳label='Smoothed Accuracy')
plt.title('Training Accuracy Improvement Trend')
plt.xlabel('Epoch')
plt.ylabel('Accuracy (%)')
plt.legend()
plt.grid(True, alpha=0.3)

# Training efficiency (accuracy per time)
plt.subplot(2, 2, 3)
efficiency = np.array(train_accuracies) / np.array(epoch_times)
plt.plot(range(1, len(efficiency) + 1), efficiency, 'g-', linewidth=2,
    ↳marker='o')
plt.title('Training Efficiency (Accuracy/Time)')
plt.xlabel('Epoch')
plt.ylabel('Accuracy % per Second')
plt.grid(True, alpha=0.3)

# Convergence analysis
plt.subplot(2, 2, 4)
loss_changes = np.abs(np.diff(train_losses))
plt.semilogy(range(2, len(train_losses) + 1), loss_changes, 'purple',
    ↳linewidth=2, marker='s')
plt.title('Loss Change Magnitude (Convergence)')
plt.xlabel('Epoch')
plt.ylabel('|Loss Change| (log scale)')
plt.grid(True, alpha=0.3)

plt.tight_layout()
plt.savefig('results/improvement_trends.png', dpi=300, bbox_inches='tight')
plt.show()

# Final performance summary
print(f"\n{'='*80}")
print("FINAL BINARY NEURAL NETWORK PERFORMANCE SUMMARY")
print(f"{'='*80}")
print(f"Model: Binary Neural Network for Plant Disease Classification")

```

```

print(f"Dataset: {len(class_names)} classes - {' '.join(class_names)}")
print(f"Total Images: {len(X)} (Train: {len(X_train)}, Test: {len(X_test)})")
print(f"Image Size: 224x224 RGB")
print(f"Model Parameters: {sum(p.numel() for p in model.parameters()),}")
print(f"Training Epochs: 20")
print(f"Training Time: {sum(epoch_times):.2f} seconds")
print(f"")
print(f"FINAL RESULTS:")
print(f"  Test Accuracy: {detailed_metrics['test_accuracy']:.2f}%")
print(f"  Test Loss: {detailed_metrics['test_loss']:.4f}")
print(f"  Average Precision: {np.mean(detailed_metrics['per_class_precision']):.3f}")
print(f"  Average Recall: {np.mean(detailed_metrics['per_class_recall']):.3f}")
print(f"  Average F1-Score: {np.mean(detailed_metrics['per_class_fscore']):.3f}")
print(f"")
print(f"MODEL COMPRESSION:")
binary_params = sum(p.numel() for name, p in model.named_parameters() if
    'hidden_layer.weight' in name)
total_params = sum(p.numel() for p in model.parameters())
print(f"  Binary Parameters: {binary_params:,} ({binary_params/total_params:.1%} of total)")
print(f"  Theoretical Storage Reduction: ~{binary_params * 32 / total_params:.1f}x for binary weights")
print(f"{'='*80}")
print("All results, graphs, and CSV files have been saved to the 'results/'
    directory.")
print(f"{'='*80}")

```

Saving model and creating additional analyses...

Model saved to: results/bnn\_plant\_disease\_model\_20250629\_131745.pth

Model info saved to: results/model\_info\_20250629\_131745.json

```

-----
ValueError                                Traceback (most recent call last)
Cell In[16], line 39
    37 # Plot ROC curve for each class
    38 for i in range(len(class_names)):
----> 39     plt.subplot(1, 3, i+1)
    40     fpr, tpr, _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    41     roc_auc = auc(fpr, tpr)

File ~/coding/drone-crop/BNN/binary nerural network/.venv/lib/python3.11/
site-packages/matplotlib/pyplot.py:1552, in subplot(*args, **kwargs)
    1549 fig = gcf()
    1551 # First, search for an existing subplot with a matching spec.

```

```

-> 1552 key = SubplotSpec._from_subplot_args(fig, args)
    1554 for ax in fig.axes:
    1555     # If we found an Axes at the position, we can reuse it if the user
    ↪passed no
    1556     # kwargs or if the Axes class and kwargs are identical.
    1557     if (ax.get_subplotspec() == key
    1558         and (kwargs == {}
    1559             or (ax._projection_init
    1560                 == fig._process_projection_requirements(**kwargs))))):

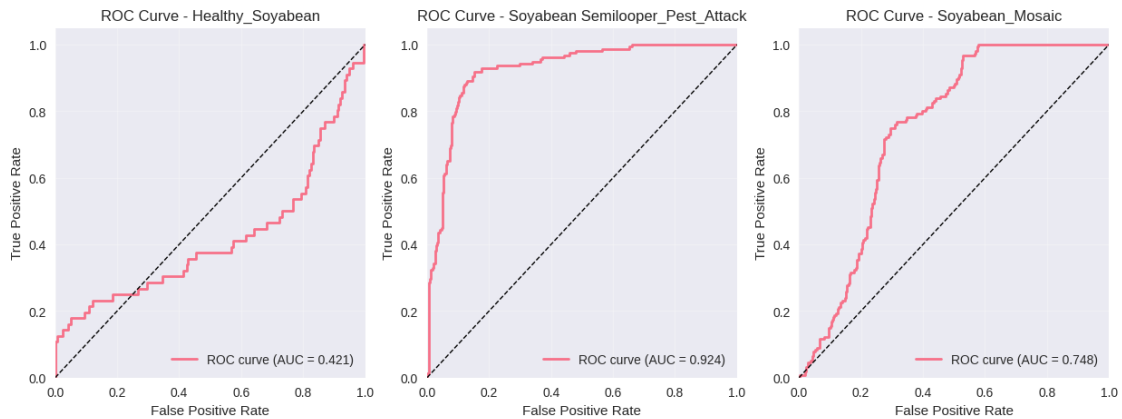
```

```

File ~/coding/drone-crop/BNN/binary nerural network/.venv/lib/python3.11/
↪site-packages/matplotlib/gridspec.py:589, in SubplotSpec.
↪_from_subplot_args(figure, args)
    587 else:
    588     if not isinstance(num, Integral) or num < 1 or num > rows*cols:
--> 589         raise ValueError(
    590             f"num must be an integer with 1 <= num <= {rows*cols}, "
    591             f"not {num!r}"
    592         )
    593     i = j = num
    594 return gs[i-1:j]

```

**ValueError: num must be an integer with 1 <= num <= 3, not 4**



```

[ ]: # Model Analysis and Binary Weights Visualization
def analyze_binary_weights(model):
    """Analyze the binary weights in the model"""
    print("Binary Weights Analysis:")
    print("=" * 50)

    # Analyze hidden layer binary weights
    with torch.no_grad():

```



```

hidden_weights = model.hidden_layer.weight
binary_weights = torch.sign(hidden_weights)

print(f"Hidden layer weights shape: {hidden_weights.shape}")
print(f"Original weights - Mean: {hidden_weights.mean():.4f}, Std: {
↳hidden_weights.std():.4f}")
print(f"Binary weights - Unique values: {torch.unique(binary_weights)}")
print(f"Binary weights - Distribution: +1: {(binary_weights == 1).sum().
↳item()}, -1: {(binary_weights == -1).sum().item()}")

# Plot weight distributions
plt.figure(figsize=(12, 4))

plt.subplot(1, 2, 1)
plt.hist(hidden_weights.cpu().numpy().flatten(), bins=50, alpha=0.7,
↳color='blue')
plt.title('Original Weights Distribution')
plt.xlabel('Weight Value')
plt.ylabel('Frequency')
plt.grid(True)

plt.subplot(1, 2, 2)
plt.hist(binary_weights.cpu().numpy().flatten(), bins=3, alpha=0.7,
↳color='red')
plt.title('Binary Weights Distribution')
plt.xlabel('Binary Weight Value')
plt.ylabel('Frequency')
plt.grid(True)

plt.tight_layout()
plt.show()

# Analyze the trained model
analyze_binary_weights(model)

# Calculate model size reduction
def calculate_model_compression():
    """Calculate the compression achieved by using binary weights"""

    # Count parameters in binary layers
    binary_params = sum(p.numel() for name, p in model.named_parameters()
        if 'hidden_layer.weight' in name)

    # Total parameters
    total_params = sum(p.numel() for p in model.parameters())

    print(f"\nModel Compression Analysis:")

```

```

print(f"Total parameters: {total_params:,}")
print(f"Binary parameters: {binary_params:,}")
print(f"Binary ratio: {binary_params/total_params:.2%}")

# In practice, binary weights can be stored using 1 bit vs 32 bits (float32)
# This gives approximately 32x compression for binary weights
theoretical_compression = binary_params * 32 / (total_params * 32 -
↪binary_params * 31)
print(f"Theoretical storage compression: {theoretical_compression:.1f}x")

calculate_model_compression()

```

## 1.2 Using Real Plant Disease Data

To use this BNN with real plant disease images, follow these steps:

### 1.2.1 1. Data Preparation

```

# Example for loading real plant disease dataset
from torchvision import datasets, transforms
from PIL import Image

# Define transforms for 224x224 RGB images
transform = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(mean=[0.485, 0.456, 0.406],
                          std=[0.229, 0.224, 0.225])
])

# Load your dataset (example structure)
# dataset/
#     healthy/
#     disease1/
#     disease2/

dataset = datasets.ImageFolder(root='path/to/your/dataset',
                              transform=transform)

```

### 1.2.2 2. Model Adaptation

- Adjust `num_classes` parameter based on your dataset
- Modify `hidden_size` for different model complexities
- Consider adding more binary layers for deeper networks

### 1.2.3 3. Training Tips

- Use data augmentation for better generalization
- Implement learning rate scheduling

- Add early stopping to prevent overfitting
- Consider using batch normalization before binary activations

#### **1.2.4 4. Performance Optimization**

- Experiment with different optimizers (SGD, AdamW)
- Try different initialization strategies for binary weights
- Implement gradient clipping for stable training
- Use mixed precision training for faster computation